



CS 224S / LINGUIST 285

Spoken Language Processing

Andrew Maas
Stanford University

Lecture 1: Introduction, ARPAbet, Articulatory Phonetics

April 3, Week 1

- Course introduction
- Course topics overview
 - Speech recognition
 - Dialog / conversational agents
 - Speech synthesis (Text to speech)
 - Affect extraction
- Very brief history
- Articulatory Phonetics
- Course Logistics
- ARPAbet transcription

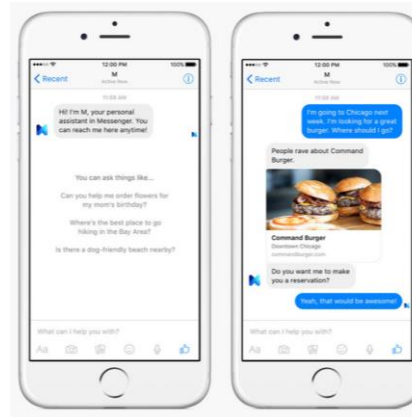
An exciting time for spoken language processing



Amazon Echo
2015




Google Home
2016




Facebook M
2015



Anki Cozmo
2016



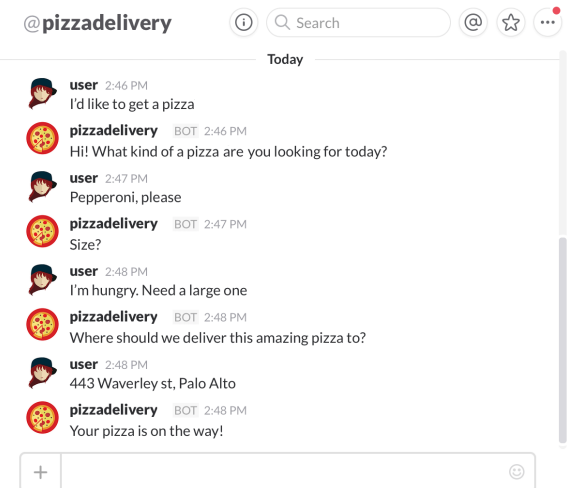

Apple
Siri
2011



Google
Assistant
2016



Microsoft
Cortana
2014



Slack Bot API
2015

LVCSR

- Large Vocabulary Continuous Speech Recognition
 - ~64,000 words
 - Speaker independent (vs. speaker-dependent)
 - Continuous speech (vs isolated-word)

Current error rates

Ballpark numbers; exact numbers depend very much on the specific corpus

Task	Vocabulary	Word Error Rate %
Digits	11	0.5
WSJ read speech	5K	1
WSJ read speech	20K	1
Broadcast news	64,000+	4
Conversational Telephone	64,000+	6

Why is conversational speech harder?



- A piece of an utterance without context



- The same utterance with more context

HSR versus ASR

Task	Vocab	ASR	Hum SR
Continuous digits	11	.5	.009
WSJ 1995 clean	5K	3	0.9
WSJ 1995 w/noise	5K	9	1.1
SWBD 2004	65K	~6	3-4?

- Conclusions:
 - Gap increases with noisy speech
 - These numbers are rough, take with grain of salt
 - We are overfitting to the benchmark datasets

HSR versus ASR

Deletions				Insertions			
SWB		CH		SWB		CH	
ASR	Human	ASR	Human	ASR	Human	ASR	Human
30: it	19: i	46: i	20: i	13: i	16: is	23: a	17: is
20: i	17: it	46: it	18: and	10: a	14: %hes	14: is	17: it
17: that	16: and	39: and	15: it	7: and	12: i	11: i	16: and
16: a	14: that	32: is	15: the	7: of	11: and	10: are	14: have
14: and	14: you	26: oh	14: is	6: you	9: it	10: you	13: a
14: oh	12: is	25: a	13: not	5: do	6: do	9: the	13: that
14: you	12: the	20: to	10: a	5: the	5: have	8: have	12: i
12: %bcack	11: a	19: that	10: in	5: yeah	5: yeah	8: that	11: %hes
12: the	10: of	19: the	10: that	4: air	5: you	7: and	10: not
11: to	9: have	18: %bcack	10: to	4: in	4: are	7: it	9: oh

Table 3: Most frequent deletion and insertion errors for humans and ASR system on SWB and CH.

SWB		CH	
ASR	Human	ASR	Human
11: and / in	16: (%hes) / oh	21: was / is	28: (%hes) / oh
9: was / is	12: was / is	16: him / them	22: was / is
7: it / that	7: (i-) / %hes	15: in / and	11: (%hes) / %bcack
6: (%hes) / oh	5: (%hes) / a	8: a / the	10: bentsy / benji
6: him / them	5: (%hes) / hmm	8: and / in	10: yeah / yep
6: too / to	5: (a-) / %hes	8: is / was	9: a / the
5: (%hes) / i	5: could / can	8: two / to	8: is / was
5: then / and	5: that / it	7: the / a	7: (%hes) / a
4: (%hes) / %bcack	4: %bcack / oh	7: too / to	7: the / a
4: (%hes) / am	4: and / in	6: (%hes) / a	7: well / oh

Table 2: Most frequent substitution errors for humans and ASR system on SWB and CH.

Why accents are hard

- A word by itself



- The word in context



So is speech recognition solved?

Why study it vs just use some API?

- In the last ~5 years
 - Dramatic reduction in LVCSR error rates (16% to 6%)
 - Human level LVCSR performance on Switchboard
 - New class of recognizers (end to end neural network)
- Understanding how ASR works enables better ASR-enabled systems
 - What types of errors are easy to correct?
 - How can a downstream system make use of uncertain outputs?
 - How much would building our own improve on an API?
- Next generation of ASR challenges as systems go live on phones and in homes

Speech Recognition Design

Intuition

- Build a statistical model of the speech-to-words process
- Collect lots and lots of speech, and transcribe all the words.
- Train the model on the labeled speech
- Paradigm: Supervised Machine Learning + Search

Dialogue (= Conversational Agents)

- Personal Assistants
 - Apple SIRI
 - Microsoft Cortana
 - Google Assistant
- Design considerations
 - Synchronous or asynchronous tasks
 - Pure speech, pure text, UI hybrids
 - Functionality versus personality

Paradigms for Dialogue

- **POMDP**

- Partially-Observed Markov Decision Processes
- Reinforcement Learning to learn what action to take
- Asking a question or answering one are just actions
 - “Speech acts”

- **Simple regular expressions and slot filling**

- Pre-built frames
 - Calendar
 - Who
 - When
 - Where
- Filled by hand-built rules
 - (“on (Mon | Tue | Wed...)”)

Paradigms for Dialogue

- **POMDP**
 - Exciting Research
 - Implemented in no commercial systems
- **Simple regular expressions and slot filling**
 - State of the art used most systems
- **Reusing new search engine technology**
 - Intent recognition / semantic parsing
- **Neural network chatbots**
 - Recent research, not really dialog yet

Extraction of Social Meaning from Speech

- Detection of student uncertainty in tutoring
 - Forbes-Riley et al. (2008)
- Emotion detection (annoyance)
 - Ang et al. (2002)
- Detection of deception
 - Newman et al. (2003)
- Detection of charisma
 - Rosenberg and Hirschberg (2005)
- Speaker stress, trauma
 - Rude et al. (2004), Pennebaker and Lay (2002)

Conversational style

- Given speech and text from a conversation
- Can we tell if a speaker is
 - Awkward?
 - Flirtatious?
 - Friendly?
- Dataset:
 - 1000 4-minute “speed-dates”
 - Each subject rated their partner for these styles
 - The following segment has been lightly signal-processed:



Speaker Recognition tasks

- Speaker Recognition
 - Speaker Verification (Speaker Detection)
 - Is this speech sample from a particular speaker
Is that Jane?
 - Speaker Identification
 - Which of these speakers does this sample come from?
Who is that?
 - Related tasks: Gender ID, Language ID
Is this a woman or a man?
- Speaker Diarization
 - Segmenting a dialogue or multiparty conversation
Who spoke when?

Applications of Speaker Recognition

- Speaker Recognition:
 - Speaker verification (binary decision)
 - Voice password
 - Telephone assistant
 - Speaker identification (one of N)
 - Criminal investigation
- Diarization
 - Transcribing meetings

TTS (= Text-to-Speech) (= Speech Synthesis)

- Produce speech from a text input
- Applications:
 - Personal Assistants
 - Apple SIRI
 - Microsoft Cortana
 - Google Assistant
 - Games
 - Airport Announcements

TTS Overview

- Main Commercial Algorithm
 - Google TTS
- Collect lots of speech (5-50 hours) from one speaker, transcribe very carefully, all the syllables and phones and whatnot
- To synthesize a sentence, patch together syllables and phones from the training data.
- Parametric synthesis shows recent gains
- First end to end neural systems in 2016

History: foundational insights 1900s-1950s




- Automaton:
 - Markov 1911
 - Turing 1936
 - McCulloch-Pitts neuron (1943)
 - <http://marr.bsee.swin.edu.au/~dtl/het704/lecture10/ann/node1.html>
 - <http://diwww.epfl.ch/mantra/tutorial/english/mcpits/html/>
 - Shannon (1948) link between automata and Markov models
- Human speech processing
 - Fletcher at Bell Labs (1920's)
- Probabilistic/Information-theoretic models
 - Shannon (1948)

Speech synthesis is old!

- Pictures and some text from Hartmut Traunmüller's web site:
 - <http://www.ling.su.se/staff/hartmut/kemplne.htm>
- **Von Kempeln** 1780 b. Bratislava 1734 d. Vienna 1804
- Leather resonator manipulated by the operator to try and copy vocal tract configuration during sonorants (vowels, glides, nasals)
- Bellows provided air stream, counterweight provided inhalation
- Vibrating reed produced periodic pressure wave


Von Kempelen:

- Small whistles controlled consonants
- Rubber mouth and nose; nose had to be covered with two fingers for non-nasals
- Unvoiced sounds: mouth covered, auxiliary bellows driven by string provides puff of air



History: Early Recognition

- 1920's Radio Rex
 - Celluloid dog with iron base held within house by electromagnet against force of spring
 - Current to magnet flowed through bridge which was sensitive to energy at 500 Hz
 - 500 Hz energy caused bridge to vibrate, interrupting current, making dog spring forward
 - The sound “e” (ARPAbet [eh]) in Rex has 500 Hz component



History: early ASR systems

- 1950's: Early Speech recognizers
 - 1952: Bell Labs single-speaker digit recognizer
 - Measured energy from two bands (formants)
 - Built with analog electrical components
 - 2% error rate for single speaker, isolated digits
 - 1958: Dudley built classifier that used continuous spectrum rather than just formants
 - 1959: Denes ASR combining grammar and acoustic probability

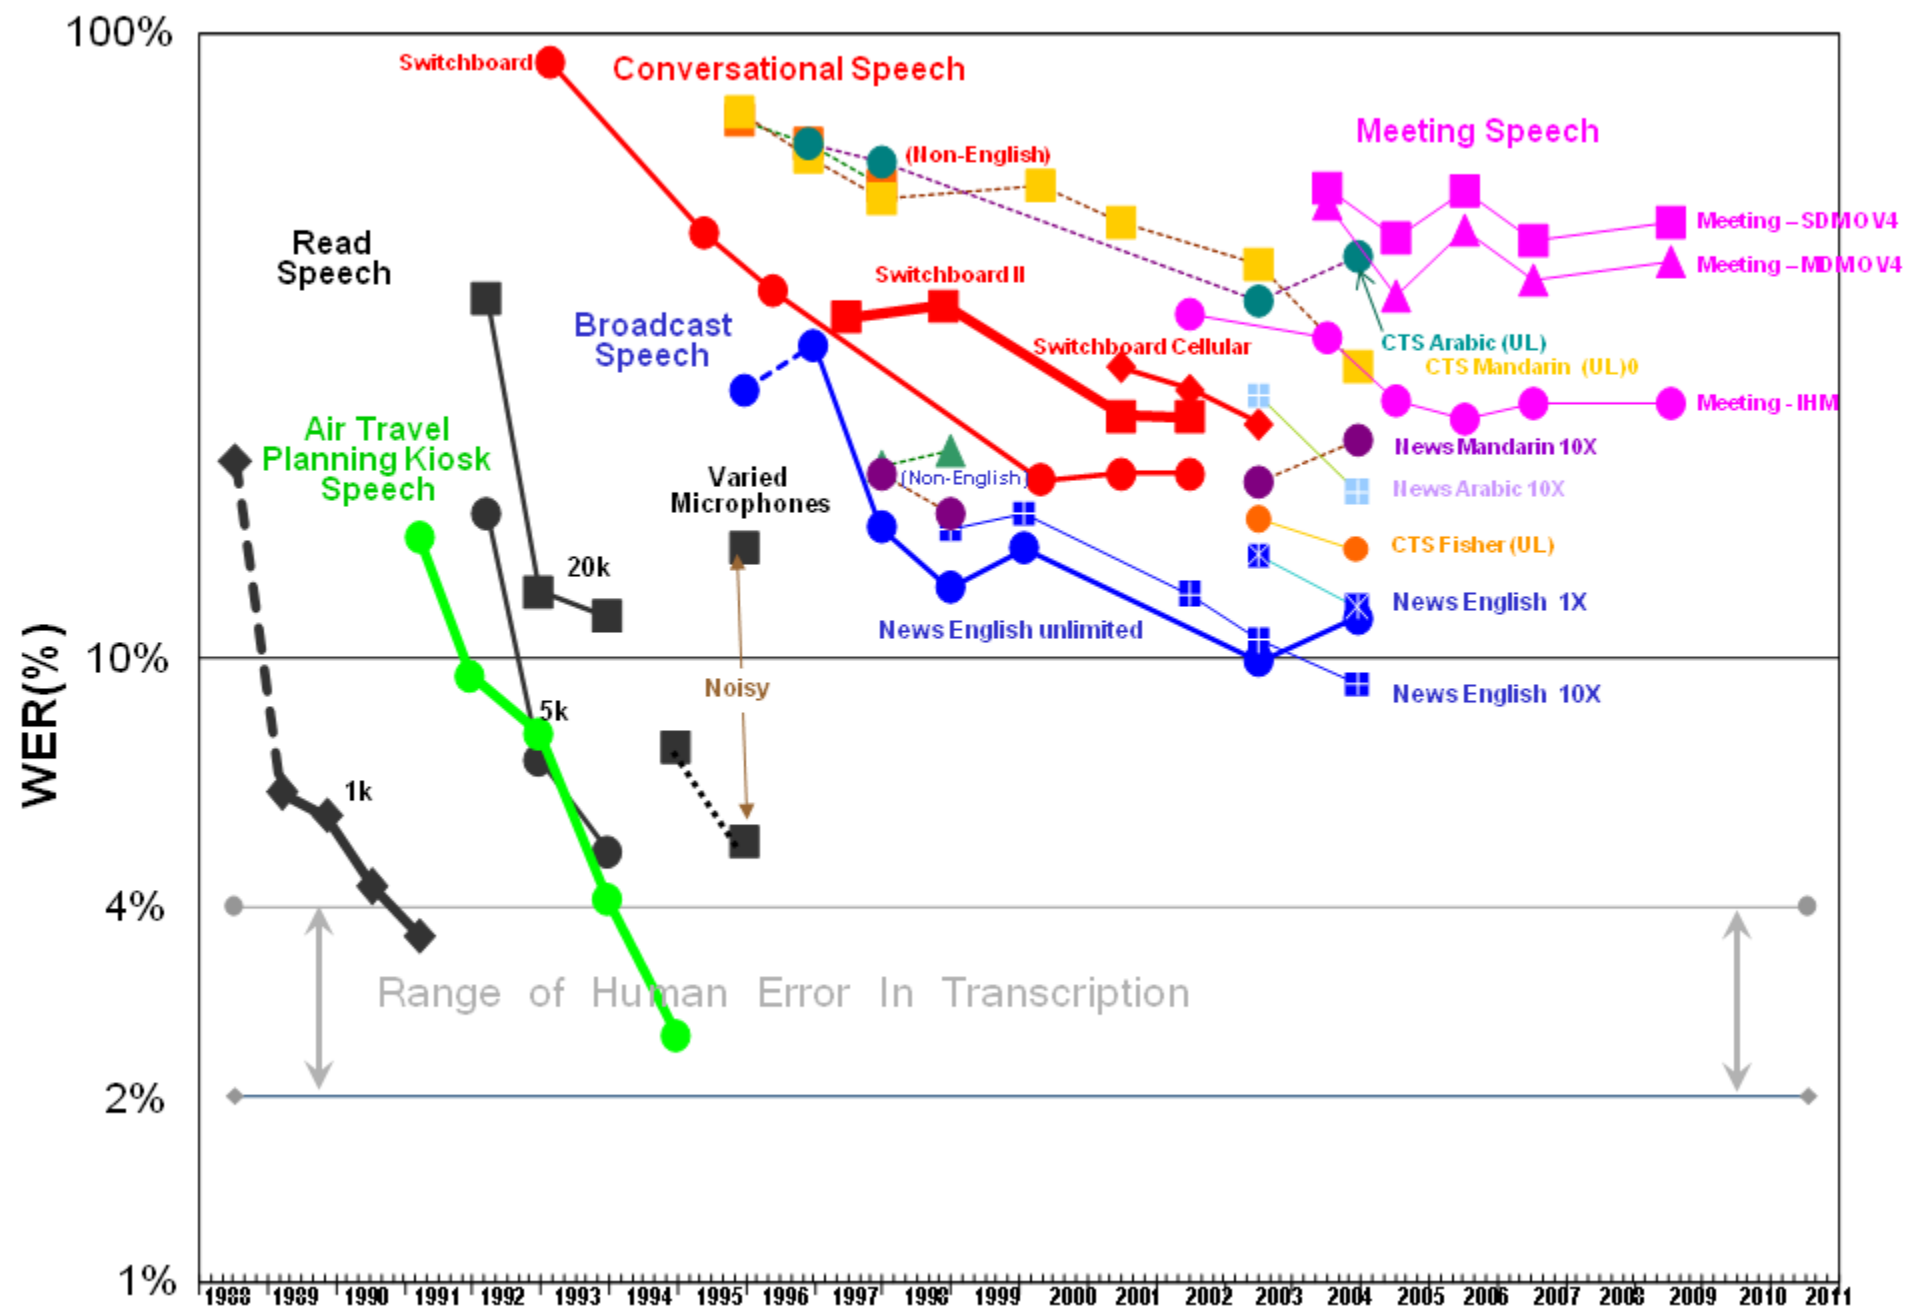
History: early ASR systems

- 1960's
 - FFT - Fast Fourier transform (Cooley and Tukey 1965)
 - LPC - linear prediction (1968)
 - 1969 John Pierce letter “Whither Speech Recognition?”
 - Random tuning of parameters,
 - Lack of scientific rigor, no evaluation metrics
 - Need to rely on higher level knowledge

ASR: 1970's and 1980's

- Hidden Markov Model 1972
 - Independent application of Baker (CMU) and Jelinek/Bahl/Mercer lab (IBM) following work of Baum and colleagues at IDA
- ARPA project 1971-1976
 - 5-year speech understanding project: 1000 word vocab, continuous speech, multi-speaker
 - SDC, CMU, BBN
 - Only 1 CMU system achieved goal
- 1980's+
 - Annual ARPA “Bakeoffs”
 - Large corpus collection
 - TIMIT
 - Resource Management
 - Wall Street Journal

NIST STT Benchmark Test History – May. '09



Course Logistics

Course Logistics

- <http://www.stanford.edu/class/cs224s>
- Homeworks released and due on Wednesdays
- Gradescope for homework submission
- Piazza for questions. Email staff only for personal/confidential questions
- Project poster session tentatively June 7 (during class time)

Admin: Requirements and Grading

- Readings:
 - Selected chapters from
 - Jurafsky & Martin. Speech and Language Processing.
 - Will mix chapters from 2nd and in progress 3rd editions
 - A few conference and journal papers
- Grading
 - Homework: 40%
 - 4 assignments. Will use Python, Tensorflow, and command line tools
 - Course Project: 50%
 - Group projects of 3 people
 - Participation: 10%

Necessary Background

- Foundations of machine learning and natural language processing
 - CS 124, CS 224N, CS 229, or equivalent experience
- Mathematical foundations of neural networks
 - Understand forward and back propagation in terms of equations
- Proficiency in Python
 - Programming heavy homeworks will use Python and Tensorflow